


latency: all outputs / primary only

(a) Cost of fetching every output



(b) GPU-resident I/O: submit vs completion

