

GPT-2 (no KV cache): per-token latency

